

cc_lead_prediction

February 11, 2022

Credit Card Lead Prediction

0.1 Problem Statement

Description: Happy Customer Bank is a mid-sized private bank that deals in all kinds of banking products, like Savings accounts, Current accounts, investment products, credit products, among other offerings.

The bank also cross-sells products to its existing customers and to do so they use different kinds of communication like tele-calling, e-mails, recommendations on net banking, mobile banking, etc.

In this case, the Happy Customer Bank wants to cross sell its credit cards to its existing customers. The bank has identified a set of customers that are eligible for taking these credit cards.

Now, the bank is looking for your help in identifying customers that could show higher intent towards a recommended credit card, given:

- Customer details (gender, age, region etc.)
- Details of his/her relationship with the bank (Channel_Code,Vintage, 'Avg_Asset_Value etc.)

0.1.1 Data Dictionary

0.1.2 Public and Private Split

Test data is further divided into Public 30% and Private 70%

0.1.3 Evaluation

The evaluation metric for this competition is **roc_auc_score** across all entries in the test set.

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
from tqdm import tqdm
from sklearn.linear_model import LogisticRegression,SGDClassifier
from sklearn.svm import LinearSVC
from sklearn.ensemble import StackingClassifier,RandomForestClassifier
from xgboost.sklearn import XGBClassifier
from sklearn.model_selection import train_test_split
```

```

from sklearn.preprocessing import OneHotEncoder,StandardScaler,Normalizer
from sklearn.calibration import CalibratedClassifierCV
from sklearn.model_selection import GridSearchCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import roc_auc_score,confusion_matrix
from sklearn.ensemble import StackingClassifier
from imblearn.over_sampling import SMOTE,RandomOverSampler
from sklearn.svm import SVC
import tensorflow as tf
from scipy.sparse import hstack
warnings.filterwarnings(action="ignore")

```

```

[3]: #loading training datasets
train_df=pd.read_csv("data/train_s3TEQDk.csv")
train_df.head()

```

```

[3]:
      ID  Gender  Age  Region_Code  Occupation  Channel_Code  Vintage  \
0  NNVBBKZB  Female   73      RG268      Other             X3        43
1  IDD62UNG  Female   30      RG277  Salaried             X1        32
2  HD3DSEMC  Female   56      RG268  Self_Employed          X3        26
3  BF3NC7KV   Male   34      RG270  Salaried             X1        19
4  TEASRWXV  Female   30      RG282  Salaried             X1        33

      Credit_Product  Avg_Account_Balance  Is_Active  Is_Lead
0                No                1045696         No         0
1                No                581988         No         0
2                No               1484315         Yes         0
3                No                470454         No         0
4                No                886787         No         0

```

```

[4]: print("List of columns::",train_df.columns.values)

```

```

List of columns:: ['ID' 'Gender' 'Age' 'Region_Code' 'Occupation' 'Channel_Code'
'Vintage'
'Credit_Product' 'Avg_Account_Balance' 'Is_Active' 'Is_Lead']

```

```

[5]: duplicated_rows=sum(train_df.duplicated(subset=['Gender', 'Age', 'Region_Code',
↪ 'Occupation', 'Channel_Code',
      'Vintage', 'Credit_Product', 'Avg_Account_Balance',
↪ 'Is_Active', 'Is_Lead']))
print("Number of duplicate rows based on sub groups of columns::
↪",duplicated_rows)

```

```

Number of duplicate rows based on sub groups of columns:: 21

```

```

[6]: # Deleting duplicates rows
train_df.drop_duplicates(subset=['Gender', 'Age', 'Region_Code', 'Occupation',
↪ 'Channel_Code',

```

```
'Vintage', 'Credit_Product', 'Avg_Account_Balance',
→ 'Is_Active', 'Is_Lead'], keep="first", inplace=True)
```

```
[7]: # training data infomation
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 245704 entries, 0 to 245724
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    245704 non-null object
1   Gender                245704 non-null object
2   Age                   245704 non-null int64
3   Region_Code           245704 non-null object
4   Occupation             245704 non-null object
5   Channel_Code           245704 non-null object
6   Vintage                245704 non-null int64
7   Credit_Product         216379 non-null object
8   Avg_Account_Balance    245704 non-null int64
9   Is_Active              245704 non-null object
10  Is_Lead                245704 non-null int64
dtypes: int64(4), object(7)
memory usage: 22.5+ MB
```

```
[8]: # Ques: How much data points are missing in every feature(percentage)
train_df.isna().sum()/train_df.shape[0]
```

```
[8]: ID                    0.000000
Gender                    0.000000
Age                       0.000000
Region_Code               0.000000
Occupation                 0.000000
Channel_Code               0.000000
Vintage                   0.000000
Credit_Product            0.119351
Avg_Account_Balance        0.000000
Is_Active                  0.000000
Is_Lead                    0.000000
dtype: float64
```

Observation * Number of data points in train datasets: 245725 * Only credit_product column has missing data points. * Around 12% of data points are missing in Credit Product feature

Ques: How can we fill this missing this value

There are many way to fill Missing value. * most common category * Predict with trained model * With Domain knowledge * Assume NaN as another category

```
[9]: #trying to find pattern how can we fill NaN values
train_df[train_df["Credit_Product"].isna()].head()
```

```
[9]:
```

	ID	Gender	Age	Region_Code	Occupation	Channel_Code	Vintage	\
6	ETQCZFEJ	Male	62	RG282	Other	X3	20	
15	UJ2NJKKL	Male	33	RG268	Self_Employed	X2	69	
31	ABPMK4WU	Female	32	RG279	Salaried	X4	15	
36	MTEIXMB9	Female	41	RG268	Self_Employed	X3	62	
40	6WX9JDVK	Female	63	RG254	Other	X3	103	

	Credit_Product	Avg_Account_Balance	Is_Active	Is_Lead
6	NaN	1056750	Yes	1
15	NaN	517063	Yes	1
31	NaN	1072850	Yes	1
36	NaN	962383	No	1
40	NaN	1249319	Yes	1

```
[10]: train_df[train_df['Credit_Product']=="Yes"].head()
```

```
[10]:
```

	ID	Gender	Age	Region_Code	Occupation	Channel_Code	Vintage	\
9	NVKTFBA2	Female	55	RG268	Self_Employed	X2	49	
11	GZ5TMYIR	Male	27	RG270	Self_Employed	X1	14	
13	KCE7JSFN	Male	31	RG254	Salaried	X1	31	
16	CNGSPYWS	Female	46	RG268	Other	X3	97	
17	VH7NBNNQ	Female	59	RG283	Other	X3	15	

	Credit_Product	Avg_Account_Balance	Is_Active	Is_Lead
9	Yes	2014239	No	0
11	Yes	502787	No	0
13	Yes	938754	No	0
16	Yes	2282502	No	1
17	Yes	2384692	No	1

```
[11]: train_df[train_df['Credit_Product']=="No"].head()
```

```
[11]:
```

	ID	Gender	Age	Region_Code	Occupation	Channel_Code	Vintage	\
0	NNVBBKZB	Female	73	RG268	Other	X3	43	
1	IDD62UNG	Female	30	RG277	Salaried	X1	32	
2	HD3DSEMC	Female	56	RG268	Self_Employed	X3	26	
3	BF3NC7KV	Male	34	RG270	Salaried	X1	19	
4	TEASRWXV	Female	30	RG282	Salaried	X1	33	

	Credit_Product	Avg_Account_Balance	Is_Active	Is_Lead
0	No	1045696	No	0
1	No	581988	No	0
2	No	1484315	Yes	0
3	No	470454	No	0

4	No	886787	No	0
---	----	--------	----	---

```
[12]: #Let see value count
train_df['Credit_Product'].value_counts(normalize=True)
```

```
[12]: No      0.667075
      Yes      0.332925
      Name: Credit_Product, dtype: float64
```

Observation:

- Above we can see that around 66% of data points have “No” as value so we can fill missing value with “No”
- or we can use some model which will train with other feature as independent variable and credit product as dependent variable
- or we can consider “NaN” as category if above idea doesn’t work well for training Models

0.2 Exploratory Data Analysis [EDA]

0.2.1 Balanced or Imbalanced Dataset

```
[13]: train_df.Is_Lead.value_counts(normalize=True)
```

```
[13]: 0      0.762771
      1      0.237229
      Name: Is_Lead, dtype: float64
```

Observation

- It’s Imbalanced Dataset so we can try SMOTE technique as Over sampling.
- Or we can try class_weight parameter in models

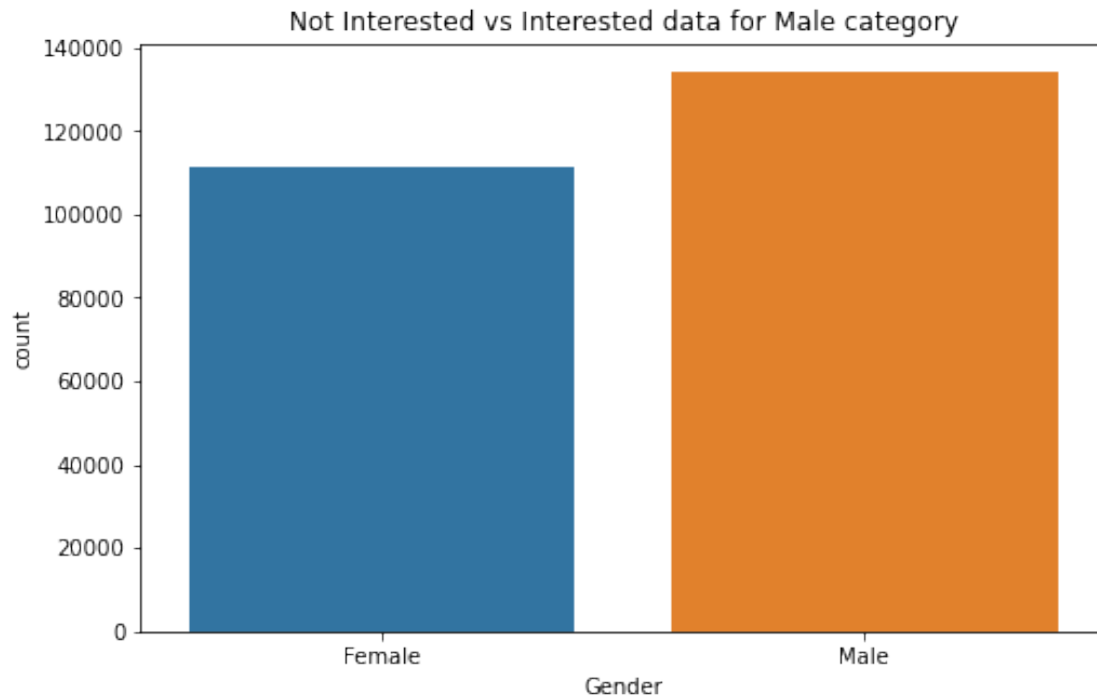
0.2.2 Gender Feature Analysis

```
[14]: # Gender Feature distribution
print("{}% of Male data present in this dataset.".
      →format(round(train_df["Gender"].value_counts(normalize=True)[0]*100)))
print("{}% of Female data present in this dataset.".
      →format(round(train_df["Gender"].value_counts(normalize=True)[1]*100)))
```

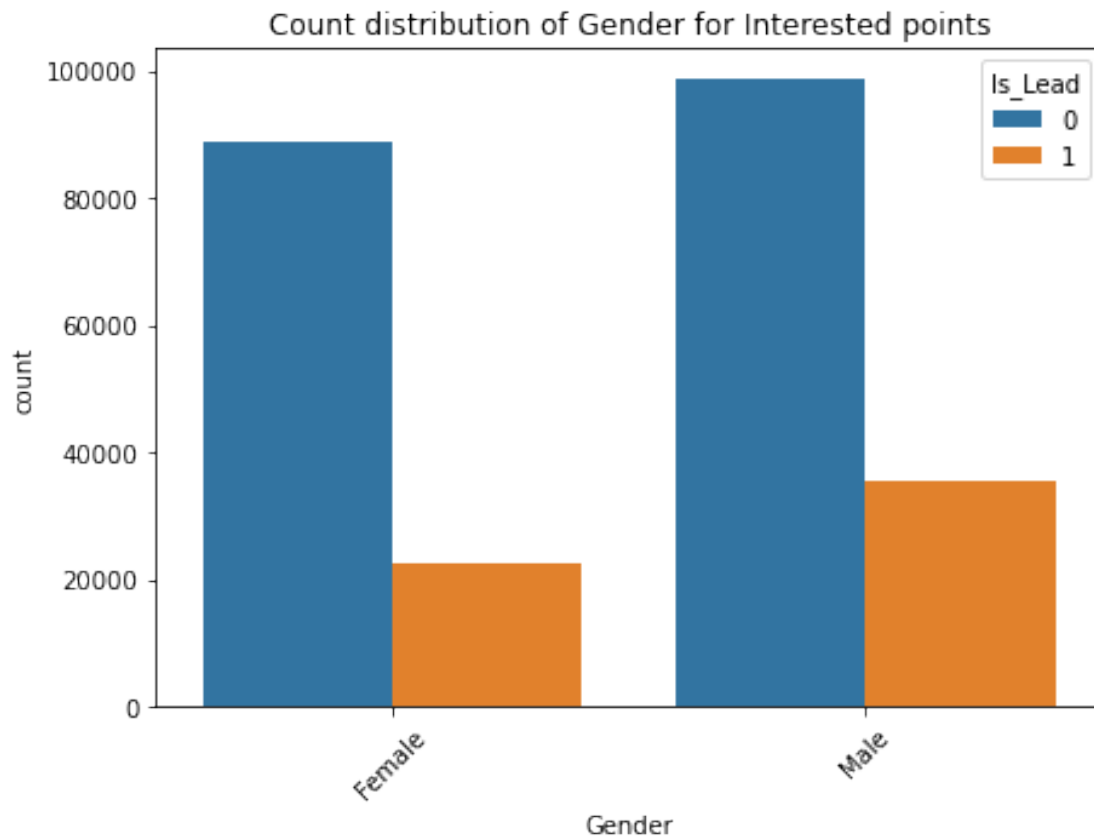
55% of Male data present in this dataset.

45% of Female data present in this dataset.

```
[15]: plt.figure(figsize=(8,5))
      plt.title("Not Interested vs Interested data for Male category")
      sns.countplot(train_df["Gender"])
      plt.show()
```




```
[16]: plt.figure(figsize=(7,5))
plt.title("Count distribution of Gender for Interested points")
sns.countplot(x=train_df["Gender"],hue=train_df["Is_Lead"])
plt.xticks(rotation=45)
plt.show()
```



Observation * this dataset has 55% of male and 45% of female data points * Female are less interested than Male

0.2.3 Age Feature Analysis

```
[17]: g=sns.FacetGrid(train_df,height=5,hue="Is_Lead").map(sns.distplot,"Age").  
      ↪add_legend()  
      g.fig.suptitle("Age Distribution")  
      plt.show()
```

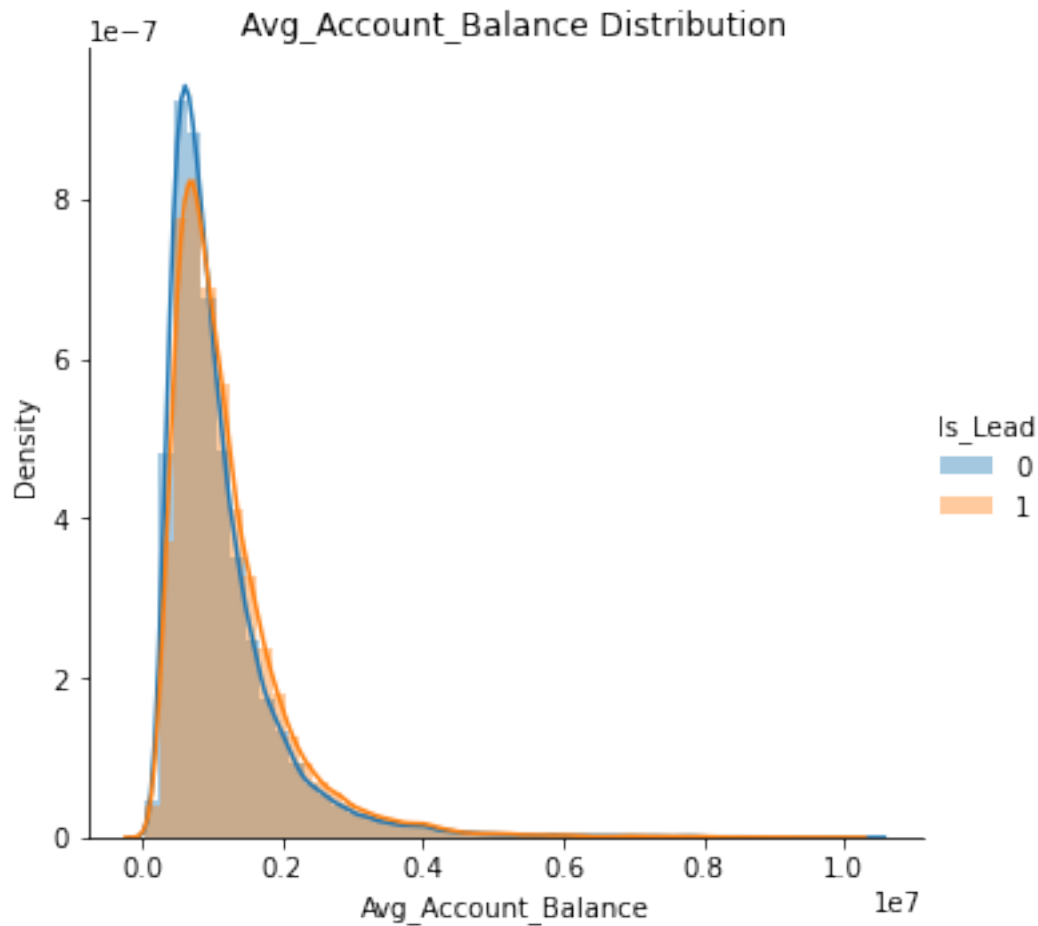


Observation: * Age distribution for “Not Interested” and “Inrerested” are overlapping. * It seems that who has age less than 35 is more “not intereseted” for credit card. * who have age >35 and age < 65 are more “Interested” for credit card. * and Age > 65 have almost equal interest for creadit card. * Based on this analogy we can convert this feature into categorical variable.

Convert to category * based on above analogy, this feature will be converted into category with below strategy. * 0-35=‘age_grp1’,35-42:‘age_grp2’,42-65:‘age_grp3’,65-above:‘age_grp4’

0.2.4 Avg_Account_Balance Feature Analysis


```
[18]: g=sns.FacetGrid(train_df,height=5,hue="Is_Lead").map(sns.
      ↳distplot,"Avg_Account_Balance").add_legend()
      g.fig.suptitle("Avg_Account_Balance Distribution")
      plt.show()
```

Observation: * distribution of Avg_Account_Balance for both Is_Lead categories are almost same. * we can see a little pick for “not interested” person who belongs between 0.0 and 0.2 range of avg acc balanced(x-axis). * so we will add this feature for training the model and will see it is helping to improve model or not.

0.2.5 Vintage Feature Analysis

```
[19]: g=sns.FacetGrid(train_df,height=5,hue="Is_Lead").map(sns.distplot,"Vintage").
      ↪add_legend()
      g.fig.suptitle("Vintage Distribution")
      plt.show()
```



Observation: * Vintage distribution for “Not Interested” and “Inrerested” are overlapping. * between vintage >0 and vintage <40 has more probability to “Not Interested” for credit card. * between vintage >40 and vintage <60 has bit more probability to “Not Interested” for credit card but alomost overlapping. * between vintage >40 and vintage <120 has more probability to “Interested” for credit card. * between vintage >120 and above has almost same probability for both Is_lead category. * Based on this analogy we can convert this feature into categorical variable.

Convert to category * based on above analogy, this feature will be converted into category with below strategy. * 0-40=‘vint_g1’,40-60=‘vint_g2’,60-80=‘vint_g3’,80-100=‘vint_g4’,100-120=‘vint_g5’,120-above=‘vint_g6’]


0.2.6 Region_Code feature

```
[20]: print("Number of Unique Region_Code :",len(train_df['Region_Code'].unique()))
```


Number of Unique Region_Code : 35

```
[21]: plt.figure(figsize=(15,5))
plt.title("Count distribution of region code for not Interested points")
```


```
sns.countplot(x=train_df[train_df["Is_Lead"]==0]["Region_Code"],color='green')
plt.xticks(rotation=45)
plt.show()
```



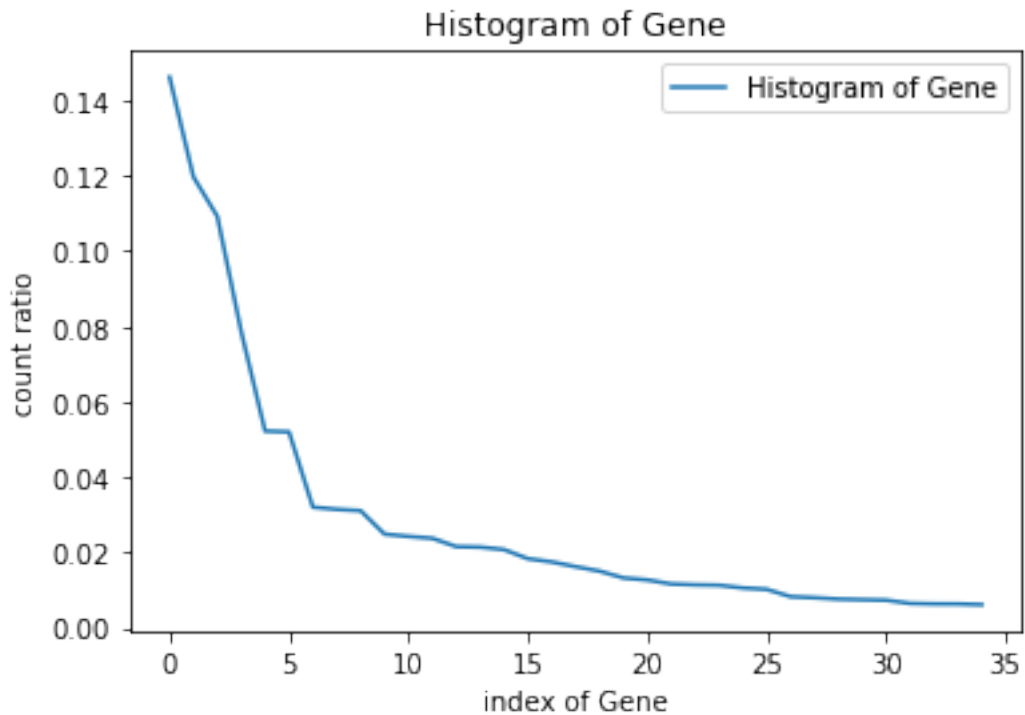
```
[22]: plt.figure(figsize=(15,5))
plt.title("Count distribution of region code for Interested points")
sns.countplot(x=train_df[train_df["Is_Lead"]==1]["Region_Code"],color='orange')
plt.xticks(rotation=45)
plt.show()
```



```
[23]: plt.figure(figsize=(15,5))
plt.title("Count distribution of region code for Interested points")
sns.countplot(x=train_df["Region_Code"],hue=train_df["Is_Lead"])
plt.xticks(rotation=45)
plt.show()
```



```
[24]: unique_gene_count=train_df['Region_Code'].value_counts(normalize=True,sort=True)
plt.plot(unique_gene_count.values,label="Histogram of Gene")
plt.title("Histogram of Gene")
plt.xlabel("index of Gene")
plt.ylabel("count ratio")
# plt.xticks([0,50,100,150,200,250])
plt.legend()
plt.show()
```



Observation: * In few region, counts are very high and the number of “not Interested” and “Interested” counts are also high in those region. * we will convert it into vector using **One Hot Encoding** and will see how much helpful for models. * And also we can trying by converting into category region group and will see how much helpful for our models.

```
[25]: #Creating reference region group category dictionary
ref_reg=train_df["Region_Code"].value_counts()
ref_reg_ind=ref_reg.index
reg_cat_dict=dict()
[reg_cat_dict.update({i:"reg_cat1"}) for i in ref_reg_ind[:5]]
[reg_cat_dict.update({i:"reg_cat2"}) for i in ref_reg_ind[5:15]]
[reg_cat_dict.update({i:"reg_cat3"}) for i in ref_reg_ind[15:25]]
[reg_cat_dict.update({i:"reg_cat4"}) for i in ref_reg_ind[25:]]
print("Tried these category but It does not work well for models")
```

Tried these category but It does not work well for models

0.2.7 Occupation Analysis

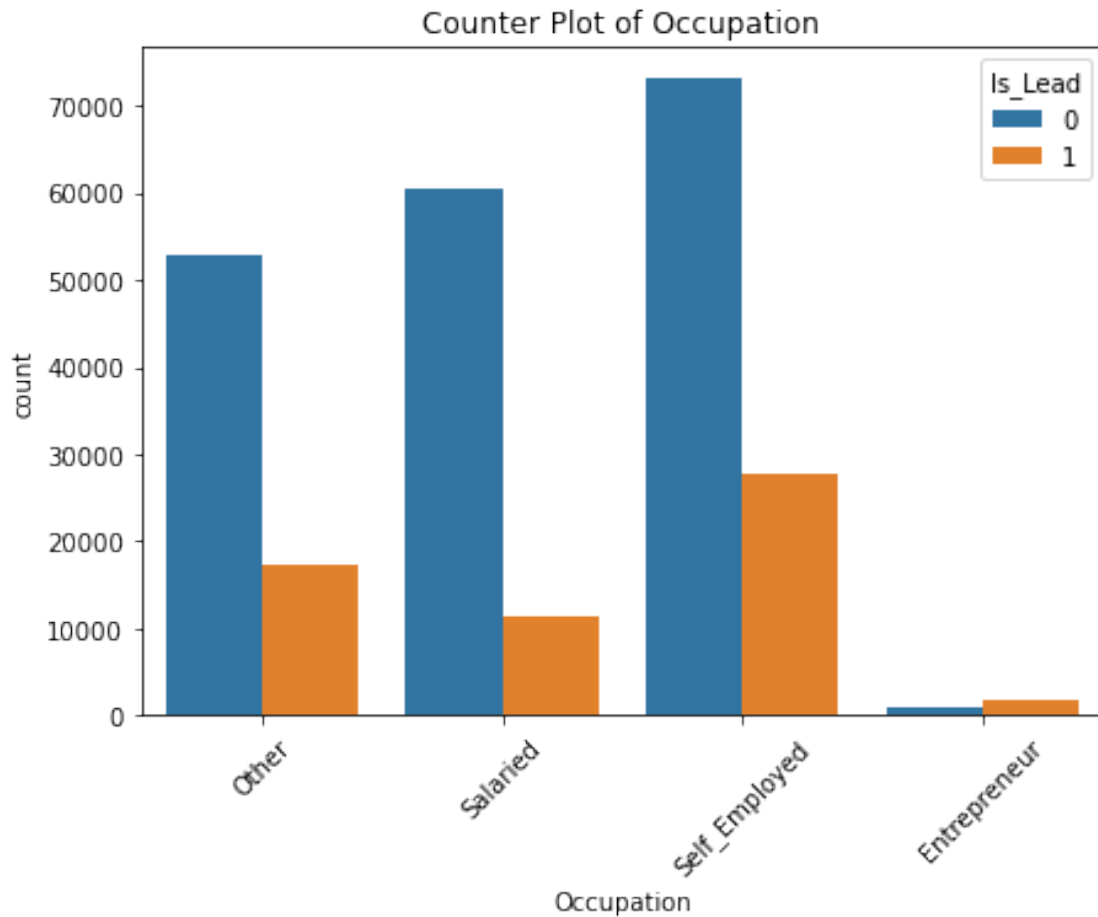
```
[26]: train_df.Occupation.value_counts(normalize=True)
```

```
[26]: Self_Employed    0.410583
      Salaried        0.292970
      Other           0.285592
      Entrepreneur    0.010855
      Name: Occupation, dtype: float64
```

```
[27]: print("Number of Unique Region_Code :",len(train_df['Occupation'].unique()))
```

Number of Unique Region_Code : 4

```
[28]: plt.figure(figsize=(7,5))
      plt.title("Counter Plot of Occupation")
      sns.countplot(x=train_df["Occupation"],hue=train_df["Is_Lead"])
      plt.xticks(rotation=45)
      plt.show()
```



Observation * Entrepreneur are more interested for credit card over not interested Entrepreneur
 * Other category are more likely to interested for credit card than salaried category

0.2.8 Channel_Code Analysis

```
[29]: print("Number of Unique Region_Code :", len(train_df['Channel_Code'].unique()))
```


Number of Unique Region_Code : 4

```
[30]: train_df.Channel_Code.value_counts(normalize=True)
```

```
[30]: X1    0.422061
      X3    0.279645
      X2    0.275628
      X4    0.022665
      Name: Channel_Code, dtype: float64
```

```
[31]: plt.figure(figsize=(7,5))
      plt.title("Counter Plot of Channel_Code")
```

```
sns.countplot(x=train_df["Channel_Code"],hue=train_df["Is_Lead"])
plt.xticks(rotation=45)
plt.show()
```



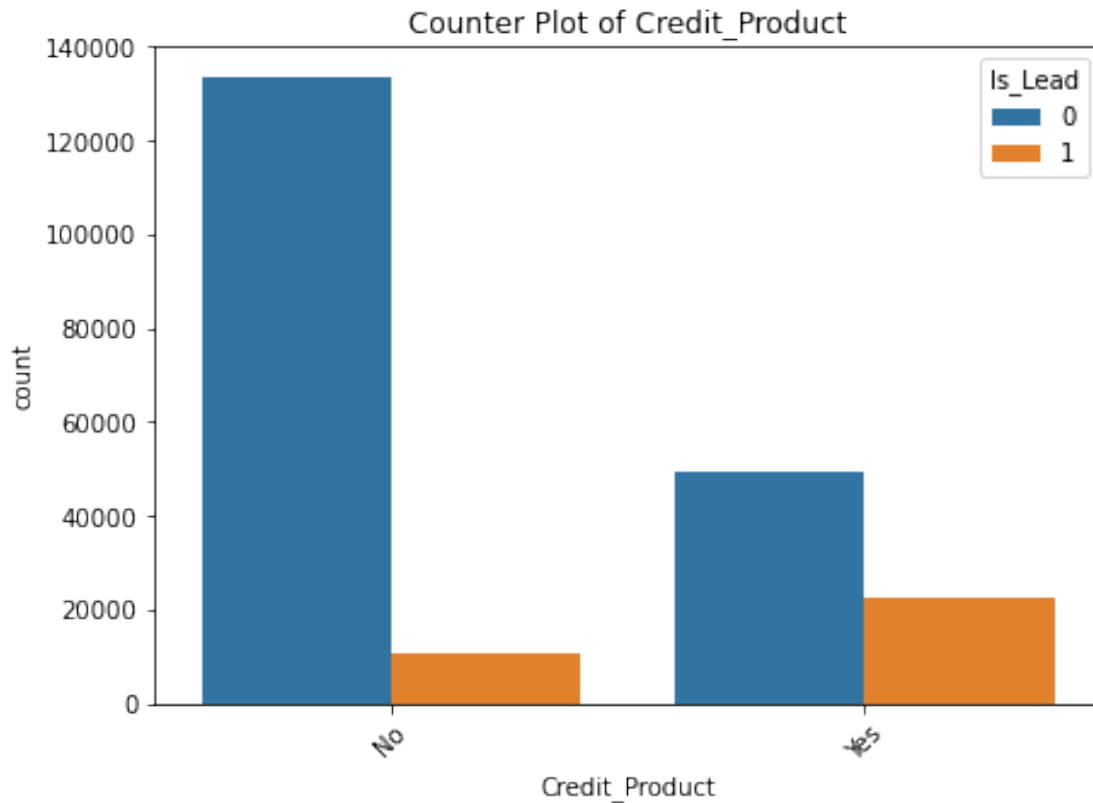
Observation * We have very less data points for category X4. It's around 2 % of whole data. * users who belongs to X1 channel are more likely to "Not interested" for credit card over "interested" X1 Channel

0.2.9 Credit_Product analysis

```
[32]: print("Number of Unique Region_Code :", len(train_df['Credit_Product'].unique()))
```

Number of Unique Region_Code : 3

```
[33]: plt.figure(figsize=(7,5))
plt.title("Counter Plot of Credit_Product")
sns.countplot(x=train_df["Credit_Product"],hue=train_df["Is_Lead"])
plt.xticks(rotation=45)
plt.show()
```



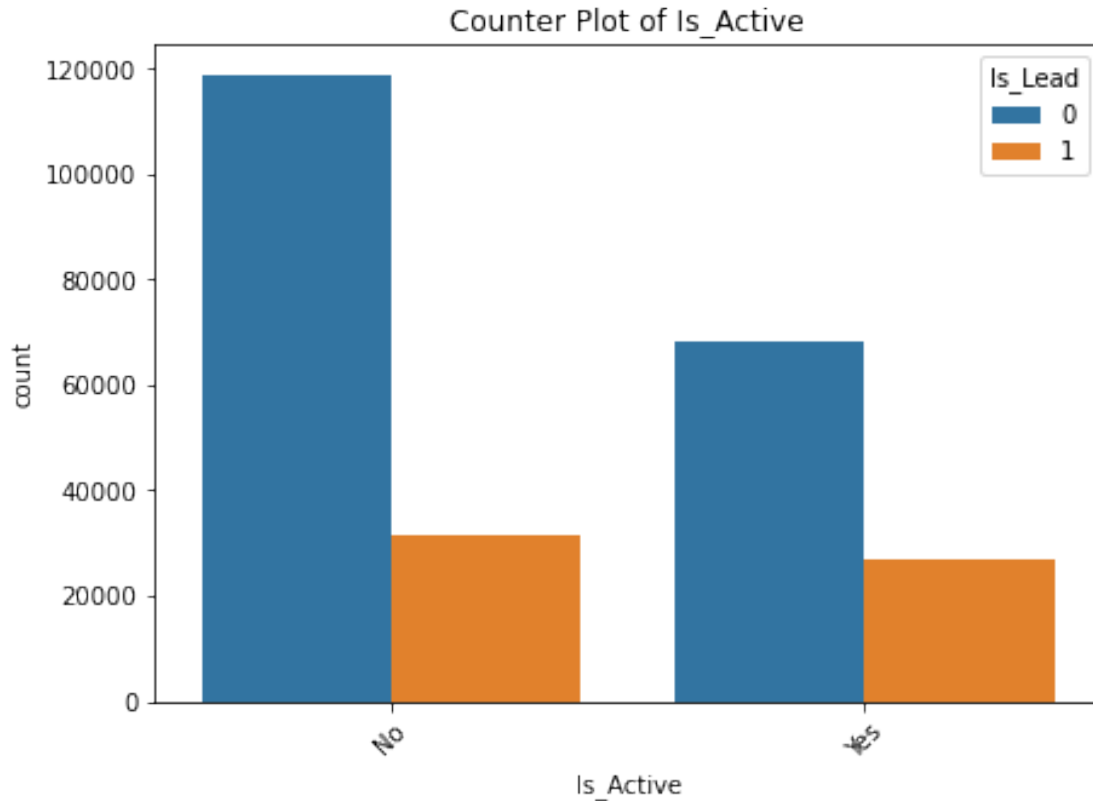
Obsevation: * There are two categories for credit product (“Yes” and “No”) along with some missing value * so for filling missing data points, we will consider third category which will be NaN.
 * user who have any active product (like Home Loan) are more likely to “Interested” for credit card rather than who have not any active product.

0.2.10 Is_Active analysis

```
[34]: print("Number of Unique Region_Code :", len(train_df['Is_Active'].unique()))
```

Number of Unique Region_Code : 2

```
[35]: plt.figure(figsize=(7,5))
plt.title("Counter Plot of Is_Active")
sns.countplot(x=train_df["Is_Active"], hue=train_df["Is_Lead"])
plt.xticks(rotation=45)
plt.show()
```

Observation: * user who is active < 3 months are more likely to Interested for credit card rather than who is not active.

0.3 Preparing Data

```
[36]: train_df[train_df.Is_Lead==1]["Credit_Product"].value_counts()
```

```
[36]: Yes    22690
      No     10623
      Name: Credit_Product, dtype: int64
```

```
[37]: train_df[train_df.Is_Lead==0]["Credit_Product"].value_counts()
```

```
[37]: No     133718
      Yes    49348
      Name: Credit_Product, dtype: int64
```

Note: * most frequent category of credit_product for not interested people is “No” * most frequent category of credit_product for Interested people is “Yes” * we can fill data point with this observation but it does not work well so I choose “NaN” as third category for all missing datapoints

```
[38]: # Defining function for preprocessing and Feature Engineering
def preprocessed_data(data):
    #Converted Age into category feature
    data["Age_group"]=pd.cut(data.
    ↪Age,bins=[0,35,45,65,120],labels=['age_grp1','age_grp2','age_grp3','age_grp4'])
    #Converted Vintage into category feature
    data["Vint_group"]=pd.cut(data.
    ↪Vintage,bins=[0,40,60,80,100,120,200],labels=['vint_g1','vint_g2','vint_g3','vint_g4','vint_g5'])
    #Concating Two feature that may help to imporve model accuracy
    data["Credit_Product"]=data["Credit_Product"].fillna("Nan")
    data["age_vintage"]=[temp[0]+"_"+temp[1] for temp in
    ↪data[["Age_group","Vint_group"]].values]
    data["age_credit"]=[temp[0]+"_"+temp[1] for temp in
    ↪data[["Age_group","Credit_Product"]].values]
    data["region_credit"]=[temp[0]+"_"+temp[1] for temp in
    ↪data[["Region_Code","Credit_Product"]].values]

    #dropping Age and Vintage after converted into category features
    data.drop(columns=["Age","Vintage"],inplace=True)
    return data
```

```
[39]: #Preprocessing training data
train_df=preprocessed_data(train_df)
train_df.head()
```

```
[39]:
```

	ID	Gender	Region_Code	Occupation	Channel_Code	Credit_Product	\
0	NNVBBKZB	Female	RG268	Other	X3	No	
1	IDD62UNG	Female	RG277	Salaried	X1	No	
2	HD3DSEMC	Female	RG268	Self_Employed	X3	No	
3	BF3NC7KV	Male	RG270	Salaried	X1	No	
4	TEASRWXV	Female	RG282	Salaried	X1	No	

	Avg_Account_Balance	Is_Active	Is_Lead	Age_group	Vint_group	\
0	1045696	No	0	age_grp4	vint_g2	
1	581988	No	0	age_grp1	vint_g1	
2	1484315	Yes	0	age_grp3	vint_g1	
3	470454	No	0	age_grp1	vint_g1	
4	886787	No	0	age_grp1	vint_g1	

	age_vintage	age_credit	region_credit
0	age_grp4_vint_g2	age_grp4_No	RG268_No
1	age_grp1_vint_g1	age_grp1_No	RG277_No
2	age_grp3_vint_g1	age_grp3_No	RG268_No
3	age_grp1_vint_g1	age_grp1_No	RG270_No
4	age_grp1_vint_g1	age_grp1_No	RG282_No

```
[40]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 245704 entries, 0 to 245724
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    245704 non-null  object
1   Gender                245704 non-null  object
2   Region_Code           245704 non-null  object
3   Occupation            245704 non-null  object
4   Channel_Code          245704 non-null  object
5   Credit_Product        245704 non-null  object
6   Avg_Account_Balance   245704 non-null  int64
7   Is_Active             245704 non-null  object
8   Is_Lead               245704 non-null  int64
9   Age_group             245704 non-null  category
10  Vint_group            245704 non-null  category
11  age_vintage           245704 non-null  object
12  age_credit            245704 non-null  object
13  region_credit         245704 non-null  object
dtypes: category(2), int64(2), object(10)
memory usage: 32.9+ MB
```

```
[41]: X=train_df.drop(columns=["Is_Lead","ID"])
      y=train_df.Is_Lead.values
```

```
[ ]:
```

0.3.1 Split into Train and CV data

```
[42]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
      ↪25, stratify=y, random_state=8)
```

```
[43]: # X_train=X
      # y_train=y
```

```
[44]: #Creating List of columns of all category features
      category_var=list(X_train.select_dtypes(include=["object","category"]).columns.
      ↪values)
```

```
[45]: #Fitting Encoder and Normalizer so that we can use it to transform train, cv
      ↪and test data
      encoder=OneHotEncoder(drop="first")
      encoder.fit(X_train[category_var])
      #Normalizer fitting
      # we can also use StandardScaler here
      norm=Normalizer()
      norm.fit(X_train["Avg_Account_Balance"].values.reshape(-1,1))
```

```
[45]: Normalizer()
```

```
[46]: ## Making list of All features
features_name=list(encoder.get_feature_names())
features_name.append("Avg_Account_Balance")
```

```
[47]: # create a function to convert category data to vector and Normalized the
      ↪ continuous variables
def encode_dataset(data):
    category_data=encoder.transform(data[category_var])
    numeric_norm=norm.transform(data["Avg_Account_Balance"].values.
    ↪ reshape(-1,1))
    # concating category data and Numerical feature
    data=hstack((category_data,numeric_norm))
    print("data shape: ",data.shape)
    return data
```

```
[48]: X_train=encode_dataset(X_train)
X_test=encode_dataset(X_test)
```

```
data shape: (184278, 189)
data shape: (61426, 189)
```

```
[49]: print("X_train shape:",X_train.shape)
print("X_test shape:",X_test.shape)
```

```
X_train shape: (184278, 189)
X_test shape: (61426, 189)
```

```
[50]: #define a function which helps to preprocess ,feature engineering and encode
      ↪ into vectors
def preprocess_and_encode_test_data(data):
    data=preprocessed_data(data)
    data=encode_dataset(data)
    return data
```

0.4 Model Training

0.4.1 Random Forest Classifier

```
[51]: #Hyper parameter tuning to RandomForestClassifier
param={"n_estimators": [100,120,150,200]}
model=RandomForestClassifier(class_weight="balanced",max_depth=10,n_jobs=-1)
# model.fit(X_train,y_train)
clf=GridSearchCV(estimator=model,param_grid=param,scoring="roc_auc",verbose=1)
clf.fit(X_train,y_train)
```

Fitting 5 folds for each of 4 candidates, totalling 20 fits

```
[51]: GridSearchCV(estimator=RandomForestClassifier(class_weight='balanced',
                                                    max_depth=10, n_jobs=-1),
                  param_grid={'n_estimators': [100, 120, 150, 200]},
                  scoring='roc_auc', verbose=1)
```

```
[52]: rf_param=clf.best_params_
print("Tunned Parameter of RF:", rf_param)
model=RandomForestClassifier(max_depth=10
                             ↵
                             ↪, n_estimators=rf_param["n_estimators"], class_weight="balanced", n_jobs=-1)
model.fit(X_train, y_train)
```

Tunned Parameter of RF: {'n_estimators': 150}

```
[52]: RandomForestClassifier(class_weight='balanced', max_depth=10, n_estimators=150,
                             n_jobs=-1)
```

Top 50 Feature (Featuere_Important)

```
[53]: ## Selecting Top 50 Features
no_of_features=50
features_dict=dict(zip(features_name, model.feature_importances_))
features_dict=sorted(features_dict.items(), key=lambda x: x[1], reverse=True)
def get_important_feature(top=10):
    features=[feat[0] for feat in features_dict[:top]]
    return features
important_feature=get_important_feature(no_of_features)
important_feature_index=[features_name.index(item) for item in ↵
                        ↪important_feature]
```

```
[54]: print("#"*10, "Top Features", "#"*10)
print(important_feature)
```

```
##### Top Features #####
['x4_No', 'x9_age_grp1_No', 'x9_age_grp3_Nan', 'x6_age_grp3', 'x7_vint_g4',
 'x9_age_grp3_No', 'x9_age_grp2_Nan', 'x4_Yes', 'x3_X3', 'x8_age_grp3_vint_g4',
 'x10_RG268_Nan', 'x9_age_grp3_Yes', 'x2_Salaried', 'x9_age_grp4_Nan', 'x3_X2',
 'x9_age_grp2_No', 'x9_age_grp1_Yes', 'x10_RG283_Nan', 'x5_Yes', 'x10_RG268_No',
 'x6_age_grp2', 'x9_age_grp2_Yes', 'x10_RG284_Nan', 'x10_RG283_No',
 'x10_RG254_Nan', 'x10_RG254_No', 'x7_vint_g3', 'x9_age_grp4_No',
 'x2_Self_Employed', 'x8_age_grp3_vint_g5', 'x8_age_grp2_vint_g1', 'x7_vint_g2',
 'x8_age_grp3_vint_g3', 'x8_age_grp4_vint_g4', 'x6_age_grp4', 'x10_RG284_No',
 'x10_RG280_Nan', 'x8_age_grp3_vint_g2', 'x8_age_grp3_vint_g1', 'x0_Male',
 'x7_vint_g5', 'x8_age_grp2_vint_g4', 'x2_Other', 'x10_RG277_Nan', 'x1_RG268',
 'x10_RG268_Yes', 'x10_RG277_No', 'x10_RG280_No', 'x8_age_grp2_vint_g3',
 'x9_age_grp4_Yes']
```

```
[55]: #Extracting top 50 features
def extract_important_features(data):
```

```

        return data.tocsr()[ :,important_feature_index]
X_train_imp=extract_important_features(X_train)
X_test_imp=extract_important_features(X_test)

```

1. RandomForestClassifier on all data

```

[56]: model=RandomForestClassifier(max_depth=10
                                     ↳
        ↳,n_estimators=rf_param["n_estimators"],class_weight="balanced",n_jobs=-1)
model.fit(X_train,y_train)

```

```

[56]: RandomForestClassifier(class_weight='balanced', max_depth=10, n_estimators=150,
                             n_jobs=-1)

```

```

[57]: #AUC on training data
pred=model.predict_proba(X_train)
auc=roc_auc_score(y_train,pred[:, 1])
print("Train AUC:",auc)

```

Train AUC: 0.8659737803098613

```

[58]: #AUC on cross validation data
pred=model.predict_proba(X_test)
auc=roc_auc_score(y_test,pred[:, 1])
print("Test AUC:",auc)

```

Test AUC: 0.8664393102039625

2. RandomForestClassifier on Top 50 features

```

[59]: model=RandomForestClassifier(max_depth=10
                                     ↳
        ↳,n_estimators=rf_param["n_estimators"],class_weight="balanced",n_jobs=-1)
model.fit(X_train_imp,y_train)

```

```

[59]: RandomForestClassifier(class_weight='balanced', max_depth=10, n_estimators=150,
                             n_jobs=-1)

```

```

[60]: #AUC on training data
pred=model.predict_proba(X_train_imp)
auc=roc_auc_score(y_train,pred[:, 1])
print("Train AUC:",auc)

```

Train AUC: 0.8708094403670058

```

[61]: #AUC on cross validation data
pred=model.predict_proba(X_test_imp)
auc=roc_auc_score(y_test,pred[:, 1])
print("Test AUC:",auc)

```

Test AUC: 0.8709047009451487

```
[ ]:
```

0.4.2 SGDClassifier

```
[62]: alpha=[10**i for i in range(-4,0)]
      param={"alpha":alpha,"penalty":["l2", 'l1', 'elasticnet']}
      model=SGDClassifier(loss="log",class_weight="balanced",learning_rate="optimal")
      #     model.fit(X_train,y_train)
      clf=GridSearchCV(estimator=model,param_grid=param,scoring="roc_auc",verbose=1)
      clf.fit(X_train,y_train)
```

Fitting 5 folds for each of 12 candidates, totalling 60 fits

```
[62]: GridSearchCV(estimator=SGDClassifier(class_weight='balanced', loss='log'),
                  param_grid={'alpha': [0.0001, 0.001, 0.01, 0.1],
                              'penalty': ['l2', 'l1', 'elasticnet']}},
                  scoring='roc_auc', verbose=1)
```

```
[63]: sgd_param=clf.best_params_
      sgd_param
```

```
[63]: {'alpha': 0.001, 'penalty': 'l2'}
```

1. SGDClassifier on All features

```
[64]: model=SGDClassifier(loss="log",class_weight="balanced",learning_rate="optimal",alpha=sgd_param)
      model.fit(X_train,y_train)
```

```
[64]: SGDClassifier(alpha=0.001, class_weight='balanced', loss='log')
```

```
[65]: #AUC on training data
      pred=model.predict_proba(X_train)
      auc=roc_auc_score(y_train,pred[:, 1])
      print("Train AUC:",auc)
```

Train AUC: 0.8645218412206594

```
[66]: #AUC on cross validation data
      pred=model.predict_proba(X_test)
      auc=roc_auc_score(y_test,pred[:, 1])
      print("Test AUC:",auc)
```

Test AUC: 0.8674847268533024

2. SGDClassifier on Top 50 features

```
[67]: model=SGDClassifier(loss="log",class_weight="balanced",learning_rate="optimal",alpha=sgd_param)
      model.fit(X_train_imp,y_train)
```

```
[67]: SGDClassifier(alpha=0.001, class_weight='balanced', loss='log')
```

```
[68]: #AUC on training data
pred=model.predict_proba(X_train_imp)
auc=roc_auc_score(y_train,pred[:, 1])
print("Train AUC:",auc)
```

Train AUC: 0.8629184442762036

```
[69]: #AUC on cross validation data
pred=model.predict_proba(X_test_imp)
auc=roc_auc_score(y_test,pred[:, 1])
print("Test AUC:",auc)
```

Test AUC: 0.8659755233259679

0.4.3 XGBOOST Classifier

```
[70]: #Hyper Tunning for XGBoostClassifier
param={"n_estimators": [100,120,150], "reg_lambda": [0.1,0.05]}
model=XGBClassifier(scale_pos_weight=3,eval_metric="logloss",learning_rate=0.
    ↪1,max_depth=5)

clf=GridSearchCV(estimator=model,param_grid=param,scoring="roc_auc",verbose=1)
clf.fit(X_train_imp,y_train)
```

Fitting 5 folds for each of 6 candidates, totalling 30 fits

```
[70]: GridSearchCV(estimator=XGBClassifier(base_score=None, booster=None,
    colsample_bylevel=None,
    colsample_bynode=None,
    colsample_bytree=None,
    enable_categorical=False,
    eval_metric='logloss', gamma=None,
    gpu_id=None, importance_type=None,
    interaction_constraints=None,
    learning_rate=0.1, max_delta_step=None,
    max_depth=5, min_child_weight=None,
    missing=nan, monotone_constraints=None,
    n_estimators=100, n_jobs=None,
    num_parallel_tree=None, predictor=None,
    random_state=None, reg_alpha=None,
    reg_lambda=None, scale_pos_weight=3,
    subsample=None, tree_method=None,
    validate_parameters=None, verbosity=None),
    param_grid={'n_estimators': [100, 120, 150],
        'reg_lambda': [0.1, 0.05]},
    scoring='roc_auc', verbose=1)
```

```
[71]: xgb_param=clf.best_params_
xgb_param
```



```
[71]: {'n_estimators': 100, 'reg_lambda': 0.1}
```

1. XGBClassifier on all Features

```
[72]: model=XGBClassifier(eval_metric="logloss",max_depth=5,reg_lambda=xgb_param["reg_lambda"]  
                        ,learning_rate=0.1,n_estimators=xgb_param["n_estimators"])  
model.fit(X_train,y_train)
```

```
[72]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
                  colsample_bynode=1, colsample_bytree=1, enable_categorical=False,  
                  eval_metric='logloss', gamma=0, gpu_id=-1, importance_type=None,  
                  interaction_constraints='', learning_rate=0.1, max_delta_step=0,  
                  max_depth=5, min_child_weight=1, missing=nan,  
                  monotone_constraints='()', n_estimators=100, n_jobs=12,  
                  num_parallel_tree=1, predictor='auto', random_state=0,  
                  reg_alpha=0, reg_lambda=0.1, scale_pos_weight=1, subsample=1,  
                  tree_method='exact', validate_parameters=1, verbosity=None)
```

```
[73]: #AUC on training data  
pred=model.predict_proba(X_train)  
auc=roc_auc_score(y_train,pred[:, 1])  
print("Train AUC:",auc)
```

Train AUC: 0.8736070816282663

```
[74]: #AUC on cross validation data  
pred=model.predict_proba(X_test)  
auc=roc_auc_score(y_test,pred[:, 1])  
print("CV AUC:",auc)
```

CV AUC: 0.87362973327029

2. XGBClassifier on Top 50 Features

```
[75]: model=XGBClassifier(eval_metric="logloss",max_depth=5,reg_lambda=xgb_param["reg_lambda"]  
                        ,learning_rate=0.1,n_estimators=xgb_param["n_estimators"])  
model.fit(X_train_imp,y_train)
```

```
[75]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
                  colsample_bynode=1, colsample_bytree=1, enable_categorical=False,  
                  eval_metric='logloss', gamma=0, gpu_id=-1, importance_type=None,  
                  interaction_constraints='', learning_rate=0.1, max_delta_step=0,  
                  max_depth=5, min_child_weight=1, missing=nan,  
                  monotone_constraints='()', n_estimators=100, n_jobs=12,  
                  num_parallel_tree=1, predictor='auto', random_state=0,  
                  reg_alpha=0, reg_lambda=0.1, scale_pos_weight=1, subsample=1,  
                  tree_method='exact', validate_parameters=1, verbosity=None)
```

```
[76]: #AUC on training data  
pred=model.predict_proba(X_train_imp)
```

```
auc=roc_auc_score(y_train,pred[:, 1])
print("Train AUC:",auc)
```

Train AUC: 0.8730313223410271

```
[77]: #AUC on cross validation data
pred=model.predict_proba(X_test_imp)
auc=roc_auc_score(y_test,pred[:, 1])
print("CV AUC:",auc)
```

CV AUC: 0.873803272448727

0.4.4 StackingClassifier

1. StackingClassifier on all features

```
[78]: # param={"n_estimators": [50,100], "max_depth": [5,10,15]}
clf1=model=XGBClassifier(eval_metric="auc",learning_rate=0.1,max_depth=5
                        ↵
                        →,n_estimators=xgb_param["n_estimators"],reg_lambda=xgb_param["reg_lambda"])
clf2=SGDClassifier(loss="log",class_weight="balanced",learning_rate="optimal",alpha=sgd_param["alpha"])

clf3=RandomForestClassifier(max_depth=10
                           ,n_estimators=100,class_weight="balanced",n_jobs=-1)

lr_clf=LogisticRegression()
model=StackingClassifier(estimators=[("xgb",clf1),("sgd",clf2),("rf",clf3)],final_estimator=lr_clf)
model.fit(X_train,y_train)
# clf=GridSearchCV(estimator=model,param_grid=param,scoring="roc_auc",verbose=1)
# clf.fit(X_train,y_train)
```

```
[78]: StackingClassifier(estimators=[('xgb',
                                     XGBClassifier(base_score=None, booster=None,
                                                    colsample_bylevel=None,
                                                    colsample_bynode=None,
                                                    colsample_bytree=None,
                                                    enable_categorical=False,
                                                    eval_metric='auc', gamma=None,
                                                    gpu_id=None, importance_type=None,
                                                    interaction_constraints=None,
                                                    learning_rate=0.1,
                                                    max_delta_step=None, max_depth=5,
                                                    min_child_weight=None,
                                                    missing=nan,
                                                    monotonic_constraints=None,
                                                    num_parallel_tree=None,
                                                    predictor=None, random_state=None,
                                                    reg_alpha=None, reg_lambda=0.1,
                                                    scale_pos_weight=None,
```

```

subsample=None, tree_method=None,
validate_parameters=None,
verbosity=None)),
('sgd',
SGDClassifier(alpha=0.001,
class_weight='balanced',
loss='log')),
('rf',
RandomForestClassifier(class_weight='balanced',
max_depth=10,
n_jobs=-1))],
final_estimator=LogisticRegression())

```

```

[79]: #AUC on training data
pred=model.predict_proba(X_train)
auc=roc_auc_score(y_train,pred[:, 1])
print("Train AUC:",auc)

```

Train AUC: 0.8724306720416939

```

[80]: #AUC on cross validation data
pred=model.predict_proba(X_test)
auc=roc_auc_score(y_test,pred[:, 1])
print("CV AUC:",auc)

```

CV AUC: 0.8729916565802007

2. StackingClassifier on Top 50 Features

```

[81]: # param={"n_estimators": [50,100], "max_depth": [5,10,15]}
clf1=model=XGBClassifier(eval_metric="auc",learning_rate=0.1,max_depth=5
    ↵
    ↪,n_estimators=xgb_param["n_estimators"],reg_lambda=xgb_param["reg_lambda"])
clf2=SGDClassifier(loss="log",class_weight="balanced",learning_rate="optimal",alpha=sgd_param["alpha"])

clf3=RandomForestClassifier(max_depth=10
    ,n_estimators=100,class_weight="balanced",n_jobs=-1)

lr_clf=LogisticRegression()
model=StackingClassifier(estimators=[("xgb",clf1),("sgd",clf2),("rf",clf3)],final_estimator=lr_clf)
model.fit(X_train_imp,y_train)
# clf=GridSearchCV(estimator=model,param_grid=param,scoring="roc_auc",verbose=1)
# clf.fit(X_train,y_train)

```

```

[81]: StackingClassifier(estimators=[('xgb',
XGBClassifier(base_score=None, booster=None,
colsample_bylevel=None,
colsample_bynode=None,
colsample_bytree=None,

```

```

enable_categorical=False,
eval_metric='auc', gamma=None,
gpu_id=None, importance_type=None,
interaction_constraints=None,
learning_rate=0.1,
max_delta_step=None, max_depth=5,
min_child_weight=None,
missing=nan,
mo...
num_parallel_tree=None,
predictor=None, random_state=None,
reg_alpha=None, reg_lambda=0.1,
scale_pos_weight=None,
subsample=None, tree_method=None,
validate_parameters=None,
verbosity=None)),
('sgd',
SGDClassifier(alpha=0.001,
class_weight='balanced',
loss='log')),
('rf',
RandomForestClassifier(class_weight='balanced',
max_depth=10,
n_jobs=-1))),
final_estimator=LogisticRegression())

```

```

[82]: #AUC on training data
pred=model.predict_proba(X_train_imp)
auc=roc_auc_score(y_train,pred[:, 1])
print("Train AUC:",auc)

```

Train AUC: 0.8729316136176766

```

[83]: #AUC on cross validation data
pred=model.predict_proba(X_test_imp)
auc=roc_auc_score(y_test,pred[:, 1])
print("Train AUC:",auc)

```

Train AUC: 0.8731234993405145

0.5 Summary

```

[97]: from tabulate import tabulate
columns=["Model","Top_50 or All","Train auc","CV auc"]
summary=[["RandomForestClassifier","All",0.8659,0.8664],
["RandomForestClassifier","Top_50",0.8708,0.8709],
["SGDClassifier","All",0.8645,0.8674],
["SGDClassifier","Top_50",0.8629,0.8659],

```

```

["XGBClassifier", "All", 0.8736, 0.8736],
["XGBClassifier", "Top_50", 0.8730, 0.8738],
["StackingClassifier", "All", 0.8724, 0.8729],
["StackingClassifier", "Top_50", 0.8729, 0.8731],
]
summary_df=pd.DataFrame(summary, columns=columns)
#https://www.geeksforgeeks.org/display-the-pandas-dataframe-in-table-style/
print(tabulate(summary_df, headers="keys", tablefmt = 'psql'))

```

	Model	Top_50 or All	Train auc	CV auc
0	RandomForestClassifier	All	0.8659	0.8664
1	RandomForestClassifier	Top_50	0.8708	0.8709
2	SGDClassifier	All	0.8645	0.8674
3	SGDClassifier	Top_50	0.8629	0.8659
4	XGBClassifier	All	0.8736	0.8736
5	XGBClassifier	Top_50	0.873	0.8738
6	StackingClassifier	All	0.8724	0.8729
7	StackingClassifier	Top_50	0.8729	0.8731

0.6 Conclusion

- Have trained many model where XGBClassifier works well
- So I have chose XGBClassifier for predicting test data

0.6.1 Best Model

```

[89]: model=XGBClassifier(eval_metric="logloss",max_depth=5,reg_lambda=xgb_param["reg_lambda"]
        ,learning_rate=0.1,n_estimators=xgb_param["n_estimators"])
model.fit(X_train,y_train)

```

```

[89]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
        colsample_bynode=1, colsample_bytree=1, enable_categorical=False,
        eval_metric='logloss', gamma=0, gpu_id=-1, importance_type=None,
        interaction_constraints='', learning_rate=0.1, max_delta_step=0,
        max_depth=5, min_child_weight=1, missing=nan,
        monotone_constraints='()', n_estimators=100, n_jobs=12,
        num_parallel_tree=1, predictor='auto', random_state=0,
        reg_alpha=0, reg_lambda=0.1, scale_pos_weight=1, subsample=1,
        tree_method='exact', validate_parameters=1, verbosity=None)

```

```

[90]: pred=model.predict_proba(X_test)
auc=roc_auc_score(y_test,pred[:, 1])
print("CV AUC:",auc)

```

CV AUC: 0.87362973327029

0.7 Creating Submission File

```
[91]: #reading Test datasets
test_data=pd.read_csv("data/test_mSzZ8RL.csv")
print("Shape of test data::",test_data.shape)
```

Shape of test data:: (105312, 10)

```
[92]: #Test data preprocessing,feature engineering and converting into vectors
test_df=preprocess_and_encode_test_data(test_data)
```

data shape: (105312, 189)

```
[93]: y_pred=model.predict_proba(test_df)[:,-1]
test_data["Is_Lead"]=y_pred
```

```
[94]: test_data=test_data[["ID", "Is_Lead"]]
```

```
[96]: test_data.to_csv("output/final_submission.csv",index=False)
```